

DEEP-CHECK

Forensic AI against deepfakes, fake identities
and document fraud

Pre-seed round · Q2 2026

Pablo López Rodríguez · Founder & CEO

pablo.lprd@gmail.com · +34 690 906 834

<https://deep-check-two.vercel.app>

Problem — Digital identity fraud is exploding

Nobody is prepared.

- **£20M stolen** from Arup in Feb 2024 through deepfake video calls.
- **+30% year-on-year** growth in synthetic identity fraud across EU banking.
- **10-15% of insurance claims** include manipulated photos across European insurers.
- **1 in 5 LinkedIn profiles** contains some fabricated credential.
- Generative AI tools (Sora 2, Flux, Midjourney) produce imagery indistinguishable from real to the human eye.

Traditional KYC, forensic experts and compliance teams are overwhelmed. The tools they use were designed before generative AI existed.

Manual verification does not scale. Adversarial generators evolve monthly. Regulation demands automated detection.

Solution — The forensic layer for the AI-fraud era

One platform. Four verification vectors:

1. **Video** — detects deepfakes, face-swaps and AI-synthesized faces (even in live video calls).
2. **Image** — detects AI-generated content from SDXL, Flux, Sora, Midjourney, DALL-E.
3. **Documents** — detects splicing, copy-move, inpainting, MRZ alteration and tampered signatures.
4. **Behavioural identity** — keystroke dynamics distinguish the legitimate person from a bot or impersonator.

Each verification is accompanied by a **cryptographically signed forensic report** (ISO/IEC 27037, eIDAS RFC 3161 timestamping) admissible as legal evidence.

Privacy-by-design: biometric analysis runs **in-browser** via WebAssembly ONNX. Biometric data never leaves the user's device.

Product — Live, not a slide deck

■ <https://deep-check-two.vercel.app>

Architecture (high level)

Layer	Component	Metric
Frontend	Next.js 16 · Vercel EU	SSR + WebAssembly ONNX
Video pipeline	V10 VideoMAE-L (304M)	AUC 0.98 · p95 98 ms
Document pipeline	DINOv2 + ELA (305M)	AUC 0.958 · p95 16 ms
Identity pipeline	Keystroke Transformer	128-D embed + bot-AUC 0.949
Semantic layer	Qwen 2.5 3B	Document reasoning
Chain of custody	ISO 27037 + RFC 3161 TSA	Append-only forensic log

Deployed in production on AWS eu-west-1 and Vercel EU. **213 000 lines of code**, 197 commits since Feb 2026. No external ML model dependencies — all nine models trained in-house.

Traction — Real revenue, validated pipeline

Metric	Value
Product live since	March 2026
Monthly revenue (3-month avg)	3,200 € · growing +50 % MoM
Signed annual contract	Yes · recurring enterprise client
Code base	213 k LOC · 197 commits · 9 branches
Trained models in production	9 (video, image, document, keystroke, semantic)

Verified pipeline in negotiation:

- **IE University** — institutional licence under internal evaluation (100 k€ ticket potential, pilot 6-15 k€). Email trail with 4 academic stakeholders CC'd.
- **Corporate Enterprise client (Zeus, sector confidential)** — technical POC advanced phase, 8-40 k€ ticket.
- **EUS Energía** — tailored identity-verification app project under discussion, 10-30 k€ project or recurring subscription.

Market — Explosive, regulated, defensible

- **Deepfake detection market:** 1.5 B\$ in 2025 → 15 B\$ in 2030 (MarketsandMarkets, CAGR 42 %).
- **Identity verification market:** 11 B\$ in 2025 → 30 B\$ in 2030.
- **Document fraud detection:** ~4 B\$ by 2028.

Regulatory tailwind (EU AI Act, Art. 50): from August 2026 onwards, companies must detect and label synthetic content. Non-compliance penalties up to **3 % of global turnover**.

Addressable segments in Europe (Year 1-3 focus):

- Insurance companies — claim fraud, ID onboarding.
- Banking and fintech — KYC, AML.
- Business schools and universities — curriculum and research.
- Media publishers — fact-checking.
- Forensic laboratories — expert-witness reports.
- Cybersecurity consultancies — resold as embedded capability.

Total SOM for Spain Y1-Y2: ~50 M€. We target 0.1 % = 500 k€.

Business Model — Multi-tier SaaS + enterprise licences

Product	Price	Target
SaaS Starter	29 €/month	Freelancers, small consultancies
SaaS Pro	79 €/month	Mid-size newsrooms, compliance teams
API Enterprise	2 000-15 000 €/year	Banks, insurers, integrators
Institutional Licence	20 000-100 000 €/year	Universities, business schools
Forensic reports	200-500 € per report	Court-bound peritaje cases
Workshops (AI Act / deepfake)	1 500-3 000 € per session	Corporate training

Gross margin: ~75 % (ML inference on reserved GPU instances, amortized over volume). **Net margin target Y3:** 35-45 %. **Revenue split Y1 projected:** 20 % SaaS + 40 % Enterprise + 25 % Institutional + 15 % services.

Competition — Why Deep-Check wins

Solution	Limitation vs Deep-Check
Microsoft Video Authenticator	Image-only. No video. No documents. No traceability.
Sensity AI	Face deepfakes only. No document forensics. No EU-compliant infra.
Truepic	Capture-time authentication (requires their camera). Cannot analyse existing content.
Reality Defender	Cloud-only API. No on-prem. No native EU privacy stance.
FacePhi	Biometric identity only. Not forensic, not multi-modal.

Defensible moat:

- Multi-vector coverage (video + image + document + behavioural + semantic) in a single platform.
- Cryptographic traceability natively — nobody else ships forensic reports that a judge can sign off on.
- Privacy-by-design: biometric inference in-browser (GDPR + AI Act aligned by default).
- EU-native infrastructure (AWS eu-west-1 + Vercel EU). No transatlantic data transfers.
- Own models, not reseller: 9 trained models, 213 k LOC of custom code. No upstream vendor lock-in.

Team & Execution

Pablo López Rodríguez, Founder & CEO

- Data / ML Engineer, 6+ years full-stack including production ML pipelines.
- Solo-built Deep-Check product, infrastructure and models in 15 months.
- Based in Madrid, full-time dedicated.

Key execution facts:

- 197 commits pushed as sole contributor.
- 9 ML models trained end-to-end.
- 2 certification processes in progress (ISO/IEC 27001, ENS-Media).
- Active fundraising + sales + product pipelines running in parallel.

Advisors / supporters identified (pending formalization):

- Fundación Don Bosco — entrepreneurship mentoring.
- Madrid Emprende — business development.
- Academic contact at IE University — product validation.

Next key hire (month 3 post-funding): 1× ML engineer (part-time) focused on iterating models against new generative adversaries.

The Ask

Raising 120 000 € pre-seed · 12 % equity · 18-month runway

Post-money valuation: 1.0 M€ (defensible given real MRR + pipeline + IP).

Concept	Amount	Why it matters
ISO/IEC 27001 + ENS-Media	25 000 €	Unlocks banking / public sector sales
Marketing B2B (insurance, banking)	20 000 €	Targeted content + paid ads + events
Part-time ML engineer (12 mo)	35 000 €	Double model-iteration velocity
Infrastructure AWS / Vercel (18 mo)	12 000 €	Reserved GPU instances, scale with clients
Legal (DPIA, terms, contracts)	8 000 €	Frictionless enterprise onboarding
Founder salary (below-market, 12 mo)	18 000 €	Sustainable dedication
Buffer / operations	2 000 €	—

Milestones with this round:

- **Month 6:** 5+ paying enterprise clients, 15 000 €+ MRR.
- **Month 12:** ISO / ENS certifications closed, 30 000 €+ MRR, AI Act enforcement live in EU.
- **Month 18:** ready for seed Series A (500 k – 1 M€) with 50 000 €+ MRR and certified client base.

Contact: pablo.lprd@gmail.com · +34 690 906 834

<https://deep-check-two.vercel.app>